

# CI AND CD AT SCALE

# SCALING JENKINS WITH DOCKER AND APACHE MESOS

Carlos Sanchez

@csanchez csanchez.org



Watch online at [carlossg.github.io/presentations](http://carlossg.github.io/presentations)

# ABOUT ME

Senior Software Engineer @ CloudBees

Contributor to the Jenkins Mesos plugin and the Java  
Marathon client

Author of Jenkins Kubernetes plugin

Long time OSS contributor at Apache, Eclipse, Puppet,...

# OUR USE CASE



Scaling Jenkins

Your mileage may vary

# SCALING JENKINS

Two options:

- More build agents per master
- More masters

# SCALING JENKINS: MORE BUILD AGENTS

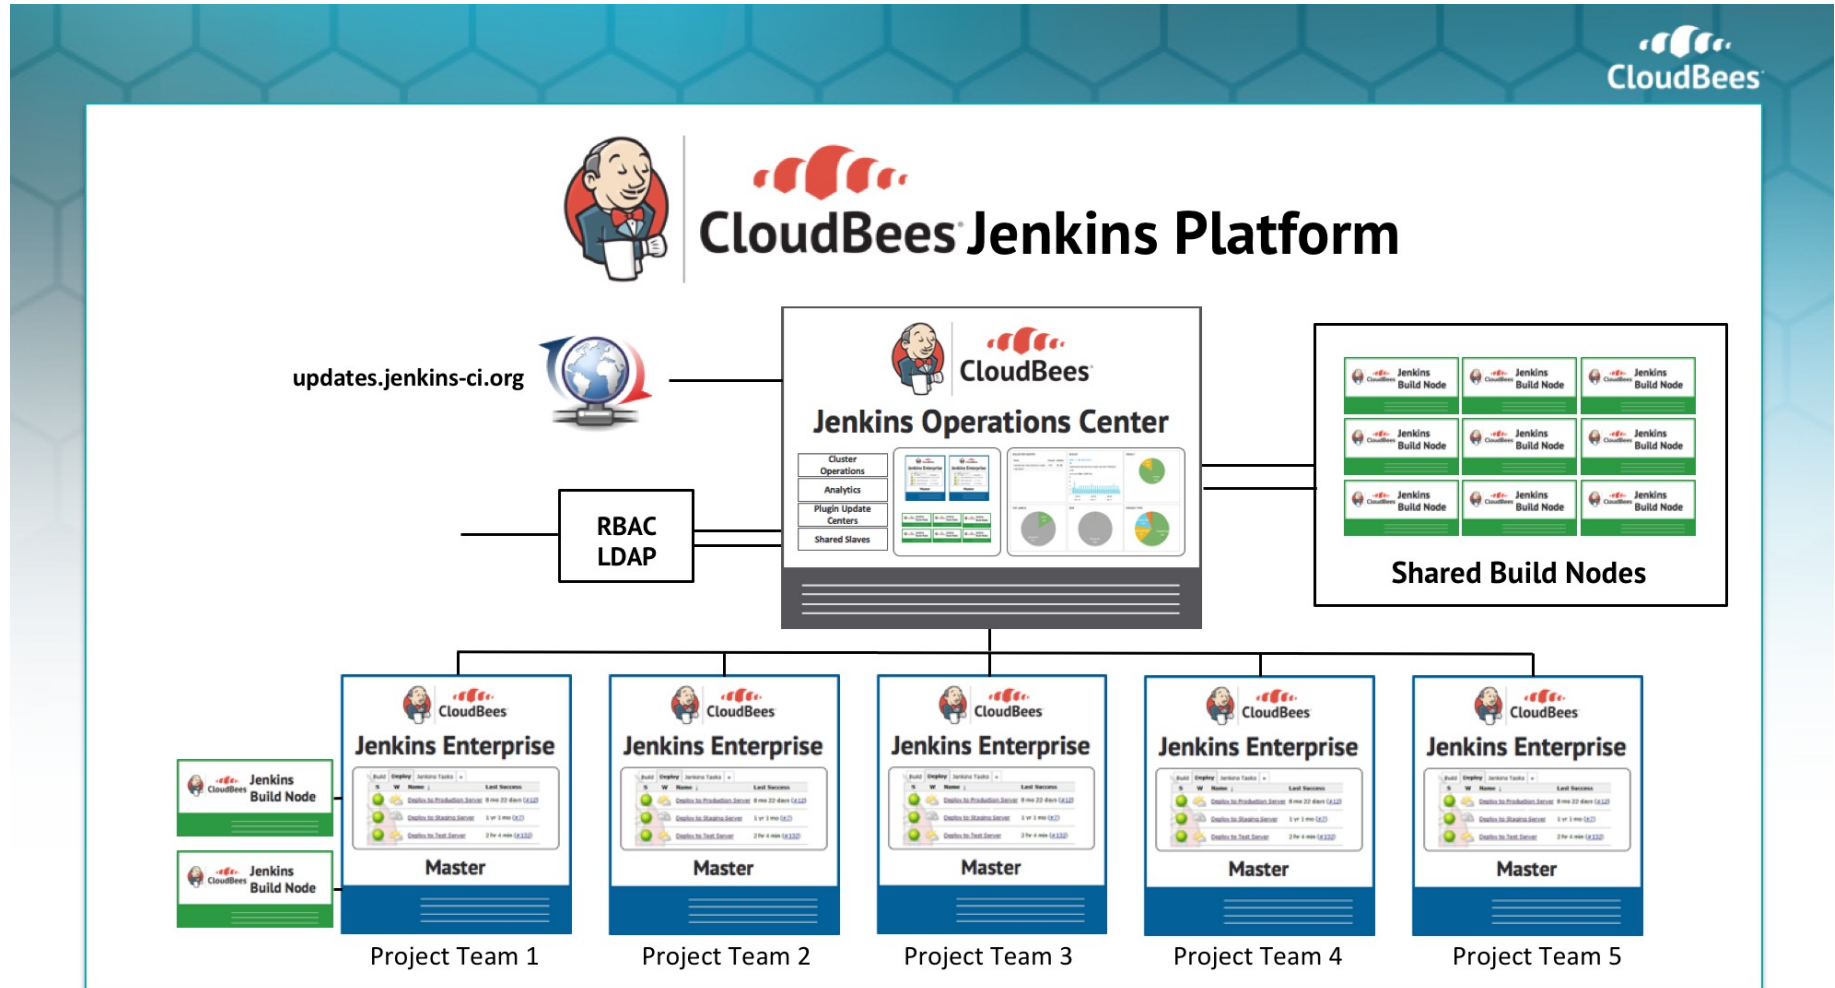
- Pros
  - Multiple plugins to add more agents, even dynamically
- Cons
  - The master is still a SPOF
  - Handling multiple configurations, plugin versions,...
  - There is a limit on how many build agents can be attached

# SCALING JENKINS: MORE MASTERS

- Pros
  - Different sub-organizations can self service and operate independently
- Cons
  - Single Sign-On
  - Centralized configuration and operation

# CLOUDBEES JENKINS ENTERPRISE EDITION

## CloudBees Jenkins Operations Center



# CLOUDBEES JENKINS PLATFORM - PRIVATE SAAS EDITION

The best of both worlds

CloudBees Jenkins Operations Center with multiple masters

Dynamic build agent creation in each master

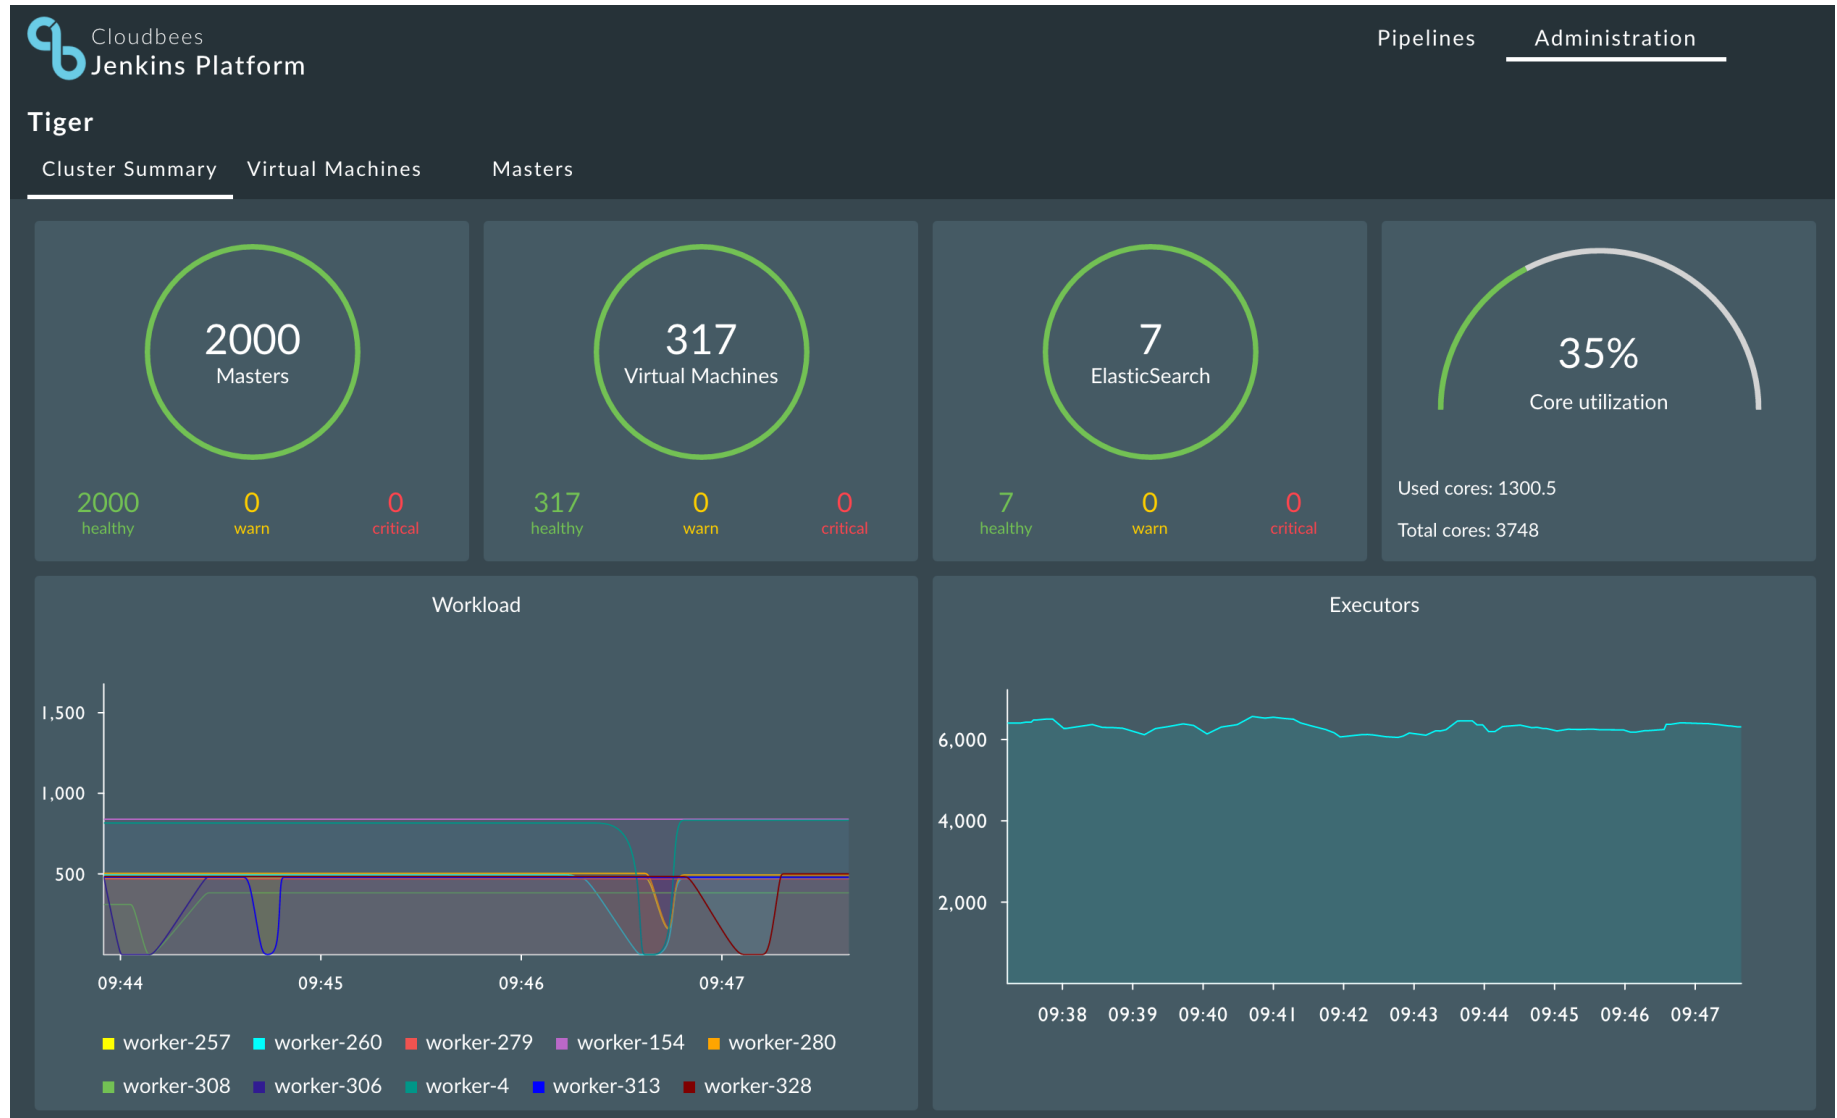
ElasticSearch for Jenkins metrics and Logstash



# BUT IT IS NOT TRIVIAL



# A 2000 JENKINS MASTERS CLUSTER




## Administration

Cluster Summary

Virtual Machines

Masters

Masters



test3-1990

2000

Master 52d6f9f7-05de-4af8-88c3-30cccbf01883

Cluster: jwpse2  
Server: 10.16.239.225:5050  
Version: 0.28.2  
Built: 3 months ago by root  
Started: yesterday  
Elected: yesterday

LOG

Slaves

Activated	313
Deactivated	0

Tasks

Staging	1,480
Starting	0
Running	11,095
Killing	0
Finished	0
Killed	2,145,109
Failed	41,123
Lost	294

Resources

	CPU's	Mem	Disk
Total	3732	12490.4 GB	32833.6 GB
Used	1500	9644.0 GB	0 B
Offered	1142.7	1965.9 GB	9537.5 GB
Idle	1089.3	880.4 GB	23296.1 GB

Active Tasks

ID	Name	State	Started ▼	Host	
test3-0942.3d13c1b2:18981c6c-61bd-456e-9bf5-995464be4327	test3-0942.3d13c1b2:18981c6c-61bd-456e-9bf5-995464be4327	STAGING		ec2-54-197-216-238.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0129.2f37ba5c:20885af1-7f5e-4458-bb5d-8b5f8a3e7aaa	test3-0129.2f37ba5c:20885af1-7f5e-4458-bb5d-8b5f8a3e7aaa	STAGING		ec2-54-197-216-238.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-1835.5714308c:b80d3dbe-6d91-4181-a04b-5e3aa83fceaeb	test3-1835.5714308c:b80d3dbe-6d91-4181-a04b-5e3aa83fceaeb	STAGING		ec2-54-226-81-206.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0702.54fa8694:b9a3cc9a-58f9-400d-b2ac-6c9dd151a963	test3-0702.54fa8694:b9a3cc9a-58f9-400d-b2ac-6c9dd151a963	STAGING		ec2-54-158-142-122.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0131.efb771db:3dffd1a19-d39d-431a-ad3e-973a4a932398	test3-0131.efb771db:3dffd1a19-d39d-431a-ad3e-973a4a932398	STAGING		ec2-54-158-164-174.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0845.f95b124b:c26c1e86-8cf1-4337-9d51-48b4b3e901f2	test3-0845.f95b124b:c26c1e86-8cf1-4337-9d51-48b4b3e901f2	STAGING		ec2-54-221-153-146.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0241.23f69555:bb19e1b7-8011-409b-9629-190ed80eca92	test3-0241.23f69555:bb19e1b7-8011-409b-9629-190ed80eca92	STAGING		ec2-52-91-32-40.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0069.3e2cd99c:6693f055-a1b6-42bb-a113-72a4c32c99ad	test3-0069.3e2cd99c:6693f055-a1b6-42bb-a113-72a4c32c99ad	STAGING		ec2-52-91-32-40.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0437.ce767edb:0a3dea36-ecf9-497f-87f3-c84d9f43756e	test3-0437.ce767edb:0a3dea36-ecf9-497f-87f3-c84d9f43756e	STAGING		ec2-52-91-88-48.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0045.5e5035ab:7e134c01-f459-443d-8dcd-2b755ae3bf84	test3-0045.5e5035ab:7e134c01-f459-443d-8dcd-2b755ae3bf84	STAGING		ec2-54-221-10-243.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-1919.d433af93:2d77536a-5eb4-4337-a0fd-b26b2a28bc84	test3-1919.d433af93:2d77536a-5eb4-4337-a0fd-b26b2a28bc84	STAGING		ec2-54-152-63-208.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0107.0baadf18:4260eb52-99c1-4453-9e49-1a011a699f47	test3-0107.0baadf18:4260eb52-99c1-4453-9e49-1a011a699f47	STAGING		ec2-54-152-63-208.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0906.d65513ff:c6f477e9-492b-4710-b1f1-c5fbbc36fa41	test3-0906.d65513ff:c6f477e9-492b-4710-b1f1-c5fbbc36fa41	STAGING		ec2-54-160-57-84.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-1495.f51d529d:b71ea06a-703f-4e12-acda-f504999f961	test3-1495.f51d529d:b71ea06a-703f-4e12-acda-f504999f961	STAGING		ec2-54-164-144-29.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-1418.8a9636b1:3923824f-d39f-4a5b-90eb-712c74e65d5c	test3-1418.8a9636b1:3923824f-d39f-4a5b-90eb-712c74e65d5c	STAGING		ec2-54-164-144-29.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-1793.700a3038:b18d3b3d-1480-4674-b4aa-ad2708a53f3c	test3-1793.700a3038:b18d3b3d-1480-4674-b4aa-ad2708a53f3c	STAGING		ec2-52-90-142-73.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0789.868c8d8b:0421730a-e875-4ddd-938e-b17b2bbe5467	test3-0789.868c8d8b:0421730a-e875-4ddd-938e-b17b2bbe5467	STAGING		ec2-54-197-213-95.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-1616.f14a1f7d:bcc9bede-40f4-4244-acf2-380c3517515f	test3-1616.f14a1f7d:bcc9bede-40f4-4244-acf2-380c3517515f	STAGING		ec2-52-91-88-48.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0799.acda253d:908732dc-10b2-4a40-8287-7b577a668f90	test3-0799.acda253d:908732dc-10b2-4a40-8287-7b577a668f90	STAGING		ec2-54-226-40-53.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-1486.cc2ccfaa:545e3b70-5fe7-41b7-bab1-31d964d1ed4e	test3-1486.cc2ccfaa:545e3b70-5fe7-41b7-bab1-31d964d1ed4e	STAGING		ec2-54-234-65-165.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0230.03416eb5:459c8c8d-a1cd-4841-ae25-537da338fe96	test3-0230.03416eb5:459c8c8d-a1cd-4841-ae25-537da338fe96	STAGING		ec2-54-234-65-165.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0324.078ea2f6:d411cb33-a481-4e7f-969c-a7a40a12818a	test3-0324.078ea2f6:d411cb33-a481-4e7f-969c-a7a40a12818a	STAGING		ec2-52-90-142-73.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-1796.88772499:d7d3da03-57ac-42df-8254-b990ed294bb8	test3-1796.88772499:d7d3da03-57ac-42df-8254-b990ed294bb8	STAGING		ec2-184-73-101-218.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0488.475b680e:2c77aa74-8da3-4710-86f7-eeb04668f7a7	test3-0488.475b680e:2c77aa74-8da3-4710-86f7-eeb04668f7a7	STAGING		ec2-54-88-19-71.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-1201.f880d741:2f9b865f-d721-4b66-a4c5-a862fbe15d10	test3-1201.f880d741:2f9b865f-d721-4b66-a4c5-a862fbe15d10	STAGING		ec2-107-22-135-75.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-0739.22d61a92:a6c84e6e-9256-40c2-9882-fa7875b91520	test3-0739.22d61a92:a6c84e6e-9256-40c2-9882-fa7875b91520	STAGING		ec2-107-22-135-75.compute-1.amazonaws.com	<a href="#">Sandbox</a>
test3-1088.58635506:20256687-36fd-4bbd-99b8-002db94601ee	test3-1088.58635506:20256687-36fd-4bbd-99b8-002db94601ee	STAGING		ec2-54-159-8-130.compute-1.amazonaws.com	<a href="#">Sandbox</a>

/masters/master-0286	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0287	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0288	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0289	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0290	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0291	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0292	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0293	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0294	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0295	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0296	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0297	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0298	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0299	2048	0.2	1 / 1	<div></div>	Running
/masters/master-0300	2048	0.2	1 / 1	<div></div>	Running


Tiger

Cluster Summary

Virtual Machines

Masters

Masters



2000

# A 2000 JENKINS MASTERS CLUSTER

- 3 Mesos masters (m3.xlarge: 4 vCPU, 15GB, 2x40 SSD)
- 317 Mesos slaves (c3.2xlarge, m3.xlarge, m4.4xlarge)
- 7 Mesos slaves dedicated to ElasticSearch: (c3.8xlarge: 32 vCPU, 60GB)

**12.5 TB - 3748 CPU**

Running 2000 masters and ~8000 concurrent jobs

# ARCHITECTURE

Docker Docker Docker



**Kernel Sanders**

@lstoll

The solution: Docker. The problem? You tell me.



Isolated Jenkins masters

Isolated build agents and jobs

Memory and CPU limits

---

*How would you design your infrastructure if  
you couldn't login? Ever.*

*Kelsey Hightower*

---

# EMBRACE FAILURE!



# CLUSTER SCHEDULING

- Running in public cloud, private cloud, VMs or bare metal
  - Starting with AWS and OpenStack
- HA and fault tolerant
- With Docker support of course

# APACHE MESOS



*A distributed systems kernel*




# ALTERNATIVES



Docker Swarm / Kubernetes

# MESOSPHERE MARATHON




For long running Jenkins masters

<1.4 does not scale with the number of apps

App definitions hit the ZooKeeper node limit

# TERRAFORM





# TERRAFORM

```
resource "aws_instance" "worker" {  
  count = 1  
  instance_type = "m3.large"  
  ami = "ami-xxxxxx"  
  key_name = "tiger-csanchez"  
  security_groups = ["sg-61bc8c18"]  
  subnet_id = "subnet-xxxxxx"  
  associate_public_ip_address = true  
  tags {  
    Name = "tiger-csanchez-worker-1"  
    "cloudbees:pse:cluster" = "tiger-csanchez"  
    "cloudbees:pse:type" = "worker"  
  }  
  root_block_device {  
    volume_size = 50  
  }  
}
```

# TERRAFORM

- State is managed
- Runs are idempotent
  - `terraform apply`
- Sometimes it is too automatic
  - Changing image id will restart all instances
- Had to fix a number of bugs, ie. retry AWS calls



**@DEVOPS\_BORAT**

DevOps Borat

To make error is human. To propagate error to all server in automatic way is **#devops**.



- Preinstall packages: Mesos, Marathon, Docker
- Cached docker images
- Other drivers: XFS, NFS,...
- Enhanced networking driver (AWS)

# MESOS FRAMEWORK

Started with Jenkins Mesos plugin

Means one framework per Jenkins master, does not scale

If master is restarted all jobs running get killed

# OUR NEW MESOS FRAMEWORK

Using Netflix Fenzo

Runs under Marathon, exposes REST API that Jenkins masters call

- Reduce number of frameworks
- Faster to spawn new build agents because framework is not started
- Pipeline durable builds, can survive a restart of the master
- Dedicated workers for builds
- Affinity

# STORAGE

Handling distributed storage

Servers can start in any host of the cluster

And they can move when they are restarted

Jenkins masters need persistent storage, agents (*typically*)  
don't

Supporting EBS (AWS) and external NFS

# **SIDEKICK CONTAINER**

A privileged container that manages mounting for other containers

Can execute commands in the host and other containers



# SIDEKICK CONTAINER *CASTLE*

Running in Marathon in each host

```
"constraints": [  
  [  
    "hostname",  
    "UNIQUE"  
  ]  
]
```

A lot of magic happening with `nsenter`  
both in host and other containers



- Jenkins master container requests data on startup using *entrypoint*
  - REST call to Castle
- Castle checks authentication
- Creates necessary storage in the backend
  - EBS volumes from snapshots
  - Directories in NFS backend

- Mounts storage in requesting container
  - EBS is mounted to host, then bind mounted into container
  - NFS is mounted directly in container
- Listens to Docker event stream for killed containers

# **CASTLE: BACKUPS AND CLEANUP**

Periodically takes snapshots from EBS volumes in AWS

Cleanups happening at different stages and periodically

**EMBRACE FAILURE!**

# PERMISSIONS

Containers should not run as root

Container user id  $\neq$  host user id

i.e. `jenkins` user in container is always 1000 but matches  
`ubuntu` user in host

# CAVEATS

Only a limited number of EBS volumes can be mounted

Docs say `/dev/sd[ f-p ]`, but `/dev/sd[ q-z ]` seem to work too

Sometimes the device gets corrupt and no more EBS volumes can be mounted there

NFS users must be centralized and match in cluster and NFS server

# MEMORY

Scheduler needs to account for container memory requirements and host available memory

Prevent containers for using more memory than allowed

Memory constrains translate to Docker `--memory`



# **WHAT DO YOU THINK HAPPENS WHEN?**

Your container goes over memory quota?



ICACHONDEO.COM

**WHAT ABOUT THE JVM?**

**WHAT ABOUT THE CHILD  
PROCESSES?**


# CPU

Scheduler needs to account for container CPU requirements  
and host available CPUs

## **WHAT DO YOU THINK HAPPENS WHEN?**

Your container tries to access more than one CPU

Your container goes over CPU limits




Totally different from memory

CPU translates into Docker `--cpu-shares`


# **OTHER CONSIDERATIONS**

# ZOMBIE REAPING PROBLEM


**Stage 1:** Nginx (PID 15) creates child process



**Stage 2:** Nginx (PID 15) exits. Its child process (PID 16) no longer has a parent and is now "orphaned"



**Stage 3:** Since PID 16 no longer has a parent, it is "adopted" by the init process, which now becomes its parent



Zombie processes are processes that have terminated but have not (yet) been waited for by their parent processes.

The init process -- PID 1 -- task is to "adopt" orphaned child processes

[source](#)



# **THIS IS A PROBLEM IN DOCKER**

Jenkins build agent run multiple processes

But Jenkins masters too, and they are long running

# TINI

Systemd or SysV init is too heavyweight for containers

---

*All Tini does is spawn a single child (Tini is meant to be run in a container), and wait for it to exit all the while reaping zombies and performing signal forwarding.*

---

## PROCESS REAPING

Docker 1.9 gave us trouble at scale, rolled back to 1.8

Lots of *defunct* processes

# NETWORKING

Jenkins masters open several ports

- HTTP
- JNLP Build agent
- SSH server (Jenkins CLI type operations)

# NETWORKING: HTTP

We use a simple `nginx` reverse proxy for

- Mesos
- Marathon
- ElasticSearch
- CJOC
- Jenkins masters

Gets destination host and port from Marathon

# NETWORKING: HTTP

Doing both

- domain based routing `master1.pse.example.com`
- path based routing `pse.example.com/master1`
  - because not everybody can touch the DNS or get a wildcard SSL certificate

# NETWORKING: JNLP

Build agents started dynamically in Mesos cluster can connect to masters internally

Build agents manually started outside cluster get host and port destination from HTTP, then connect directly

# NETWORKING: SSH

## SSH Gateway Service

Tunnel SSH requests to the correct host

Simple configuration needed in client

```
Host=*.ci.cloudbees.com  
ProxyCommand=ssh -q -p 22 ssh.ci.cloudbees.com tunnel %h
```

allows to run


```
ssh master1.ci.cloudbees.com
```

# SCALING

New and interesting problems





A close-up, low-angle shot of a man in a suit, looking down with a serious expression. The lighting is dramatic, with strong shadows on his face. The background is blurred, showing another person in a suit.

If you never used Docker in  
production, leave the room now

# TERRAFORM AWS

- Instances
- Keypairs
- Security Groups
- S3 buckets
- ELB
- VPCs

# **AWS**

Resource limits: VPCs, S3 snapshots, some instance sizes

Rate limits: affect the whole account

Retrying is your friend, but with exponential backoff

# AWS

Running with a patched Terraform to overcome timeouts  
and AWS *eventual consistency*

```
<?xml version="1.0" encoding="UTF-8"?>
<DescribeVpcsResponse xmlns="http://ec2.amazonaws.com/doc/2015-10-01/"
  <requestId>8f855bob-3421-4cff-8c36-4b517eb0456c</requestId>
  <vpcSet>
    <item>
      <vpcId>vpc-30136159</vpcId>
      <state>available</state>
      <cidrBlock>10.16.0.0/16</cidrBlock>
      ...
    </item>
  </vpcSet>
</DescribeVpcsResponse>
2016/05/18 12:55:57 [DEBUG] [aws-sdk-go] DEBUG: Response ec2/Describe
--[ RESPONSE ] -----
HTTP/1.1 400 Bad Request
<Response><Errors><Error><Code>InvalidVpcID.NotFound</Code><Message>
The vpc ID 'vpc-30136159' does not
exist</Message></Error></Errors>
```

# TERRAFORM OPENSTACK

- Instances
- Keypairs
- Security Groups
- Load Balancer
- Networks

# OPENSTACK

Custom flavors

Custom images

Different CLI commands

There are not two OpenStack installations that are the same

# GRACIAS

[csanchez.org](http://csanchez.org)



[csanchez](https://twitter.com/csanchez)

[carlossg](https://twitter.com/carlossg)

cloudbees®