

# DEDICATED INFRASTRUCTURE IN A MULTITENANT WORLD



Carlos Sanchez / [csanchez.org](https://csanchez.org) / [@csanchez](https://twitter.com/csanchez)

Cloud Engineer

Adobe Experience Manager Cloud Service

Author of Jenkins Kubernetes plugin

Long time OSS contributor at Jenkins, Apache Maven,  
Puppet,...

# **ADOBE EXPERIENCE MANAGER**

Content Management System

Digital Asset Management

Digital Enrollment and Forms

Used by many Fortune 100 companies

An existing distributed Java OSGi application

Using OSS components from Apache Software Foundation

A huge market of extension developers

**AEM ON KUBERNETES**

Running on Azure

18+ clusters and growing

Multiple regions: US, Europe, Australia, Singapore, Japan, more coming

Adobe has a dedicated team managing clusters for multiple products

Customers can run their own code

Cluster permissions are limited for security

Traffic leaving the clusters must be encrypted



Using namespaces to provide a scope

- network isolation
- quotas
- permissions

More details on our Kubernetes setup in my KubeCon 2020 talk

<https://tinyurl.com/csanchez-kcna20>



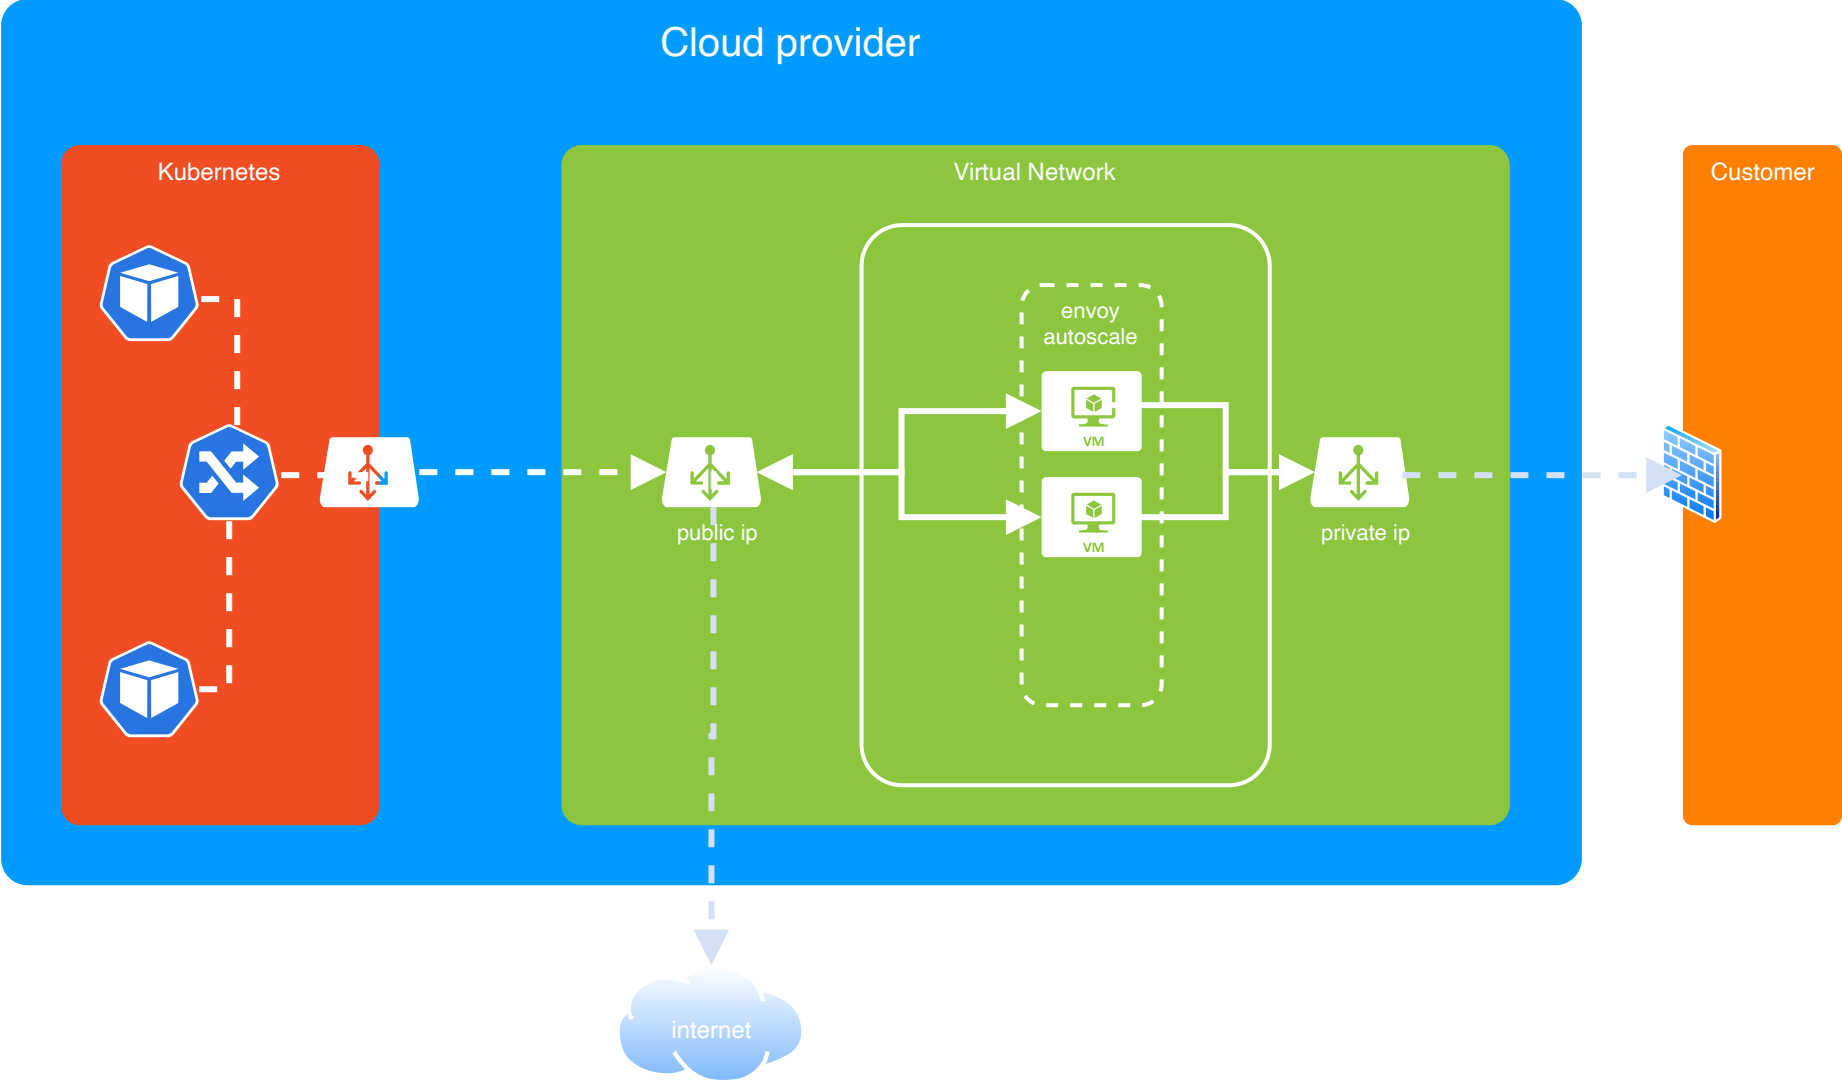
# DEDICATED INFRASTRUCTURE

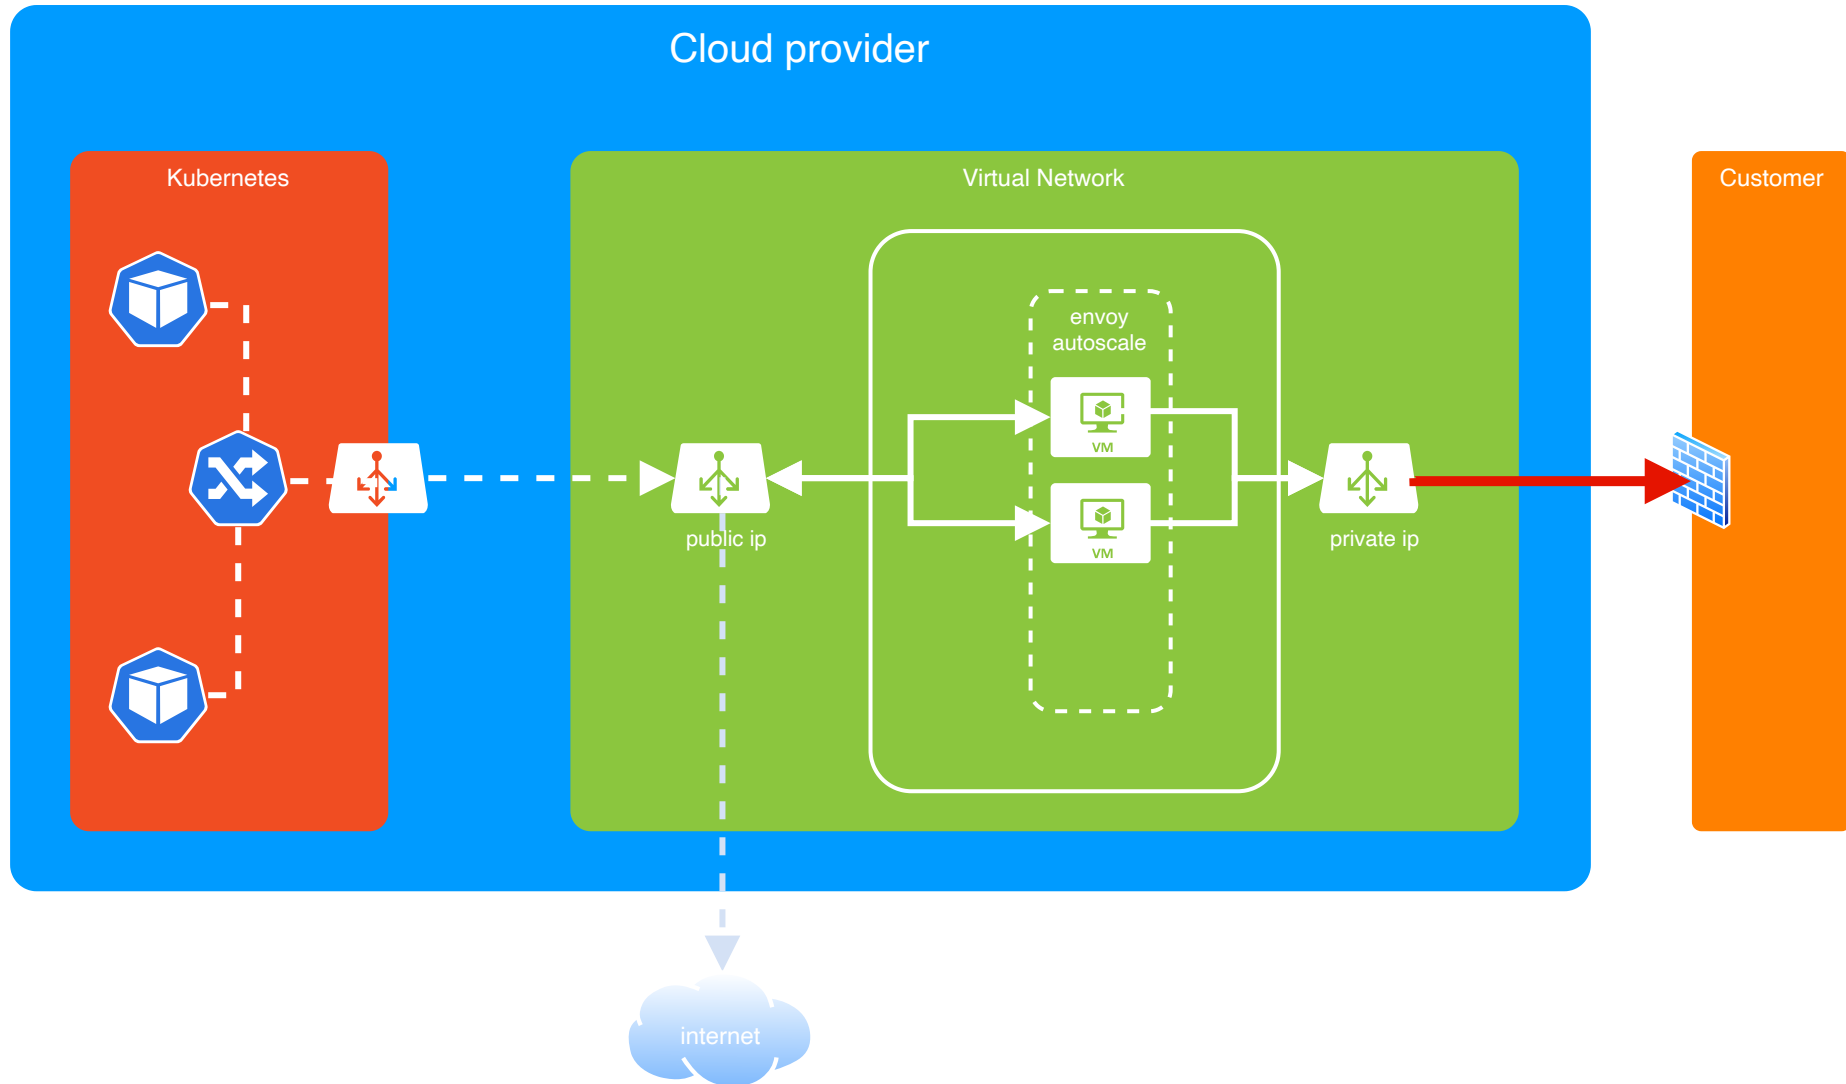
Customers want to have dedicated infrastructure

- Egress IPs
- Private connections (VNET peering, Private Link, ExpressRoute, Direct Connect,...)
- VPN

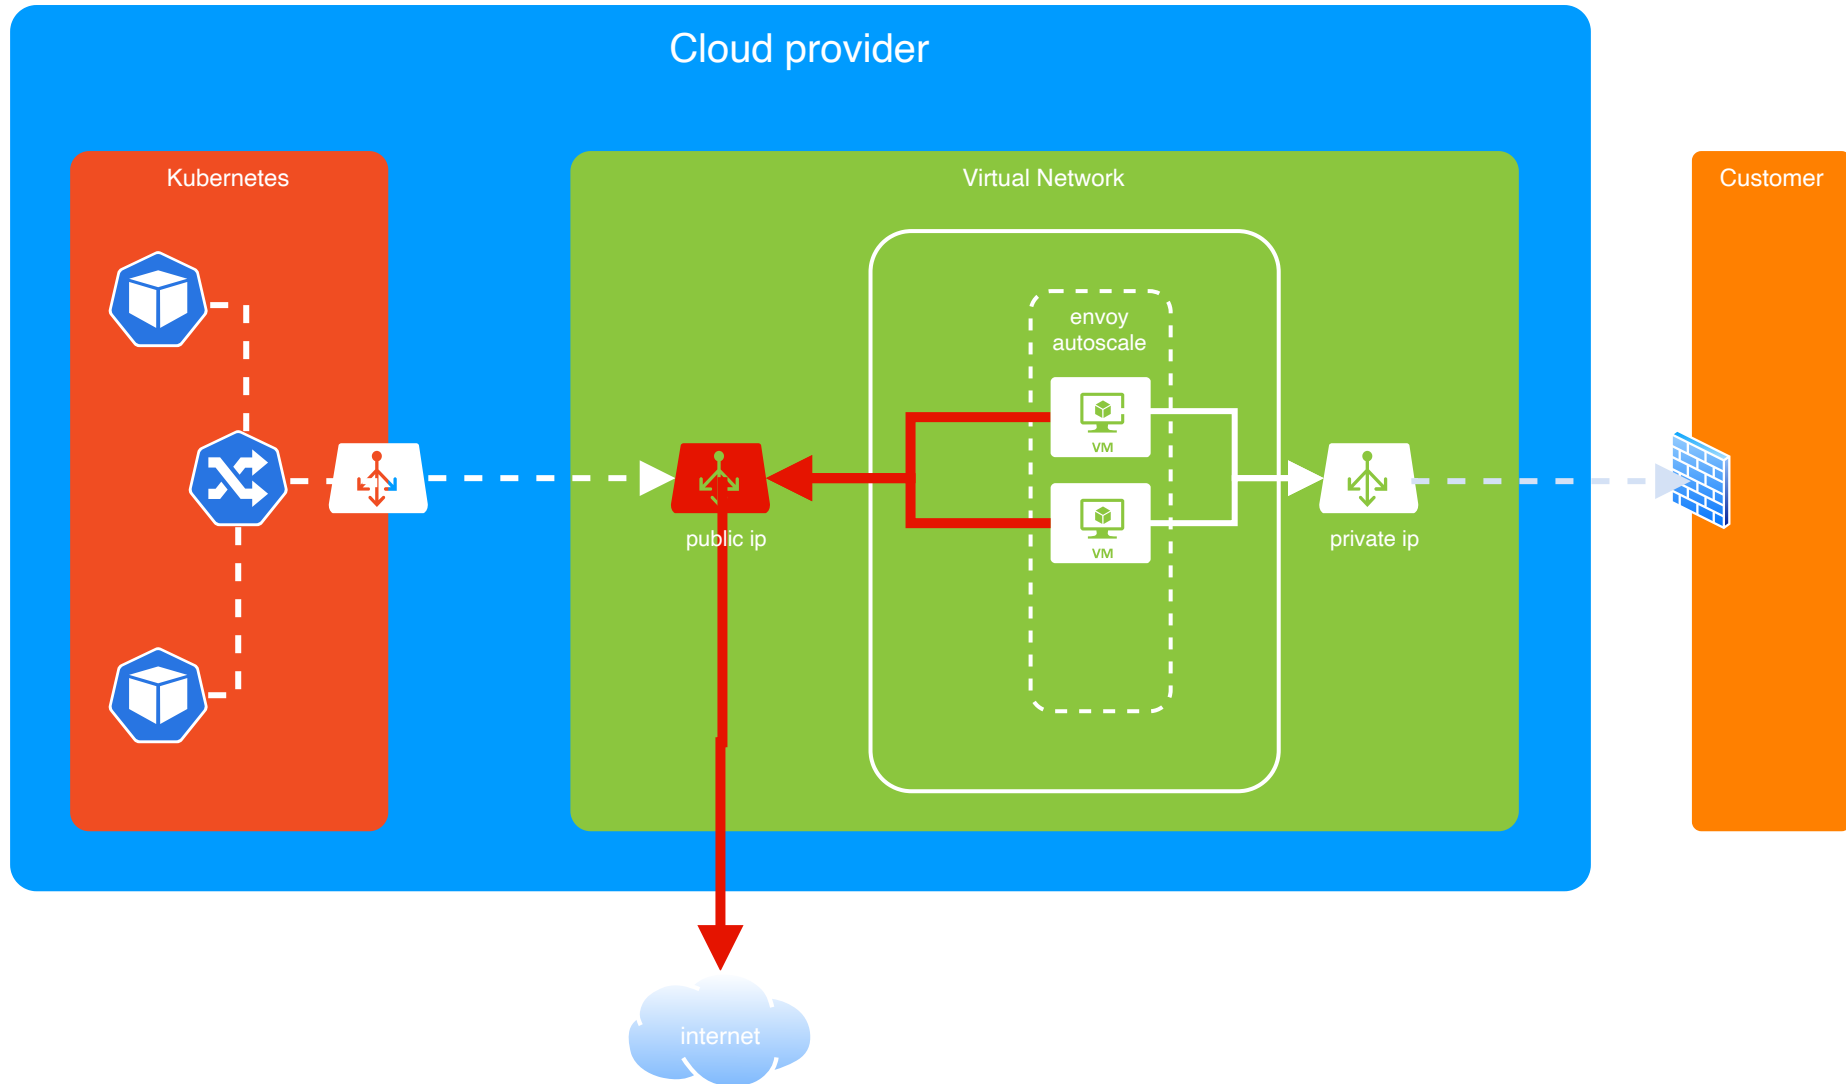
# ENVOY

Running Envoy on VMs and pod sidecars in Kubernetes

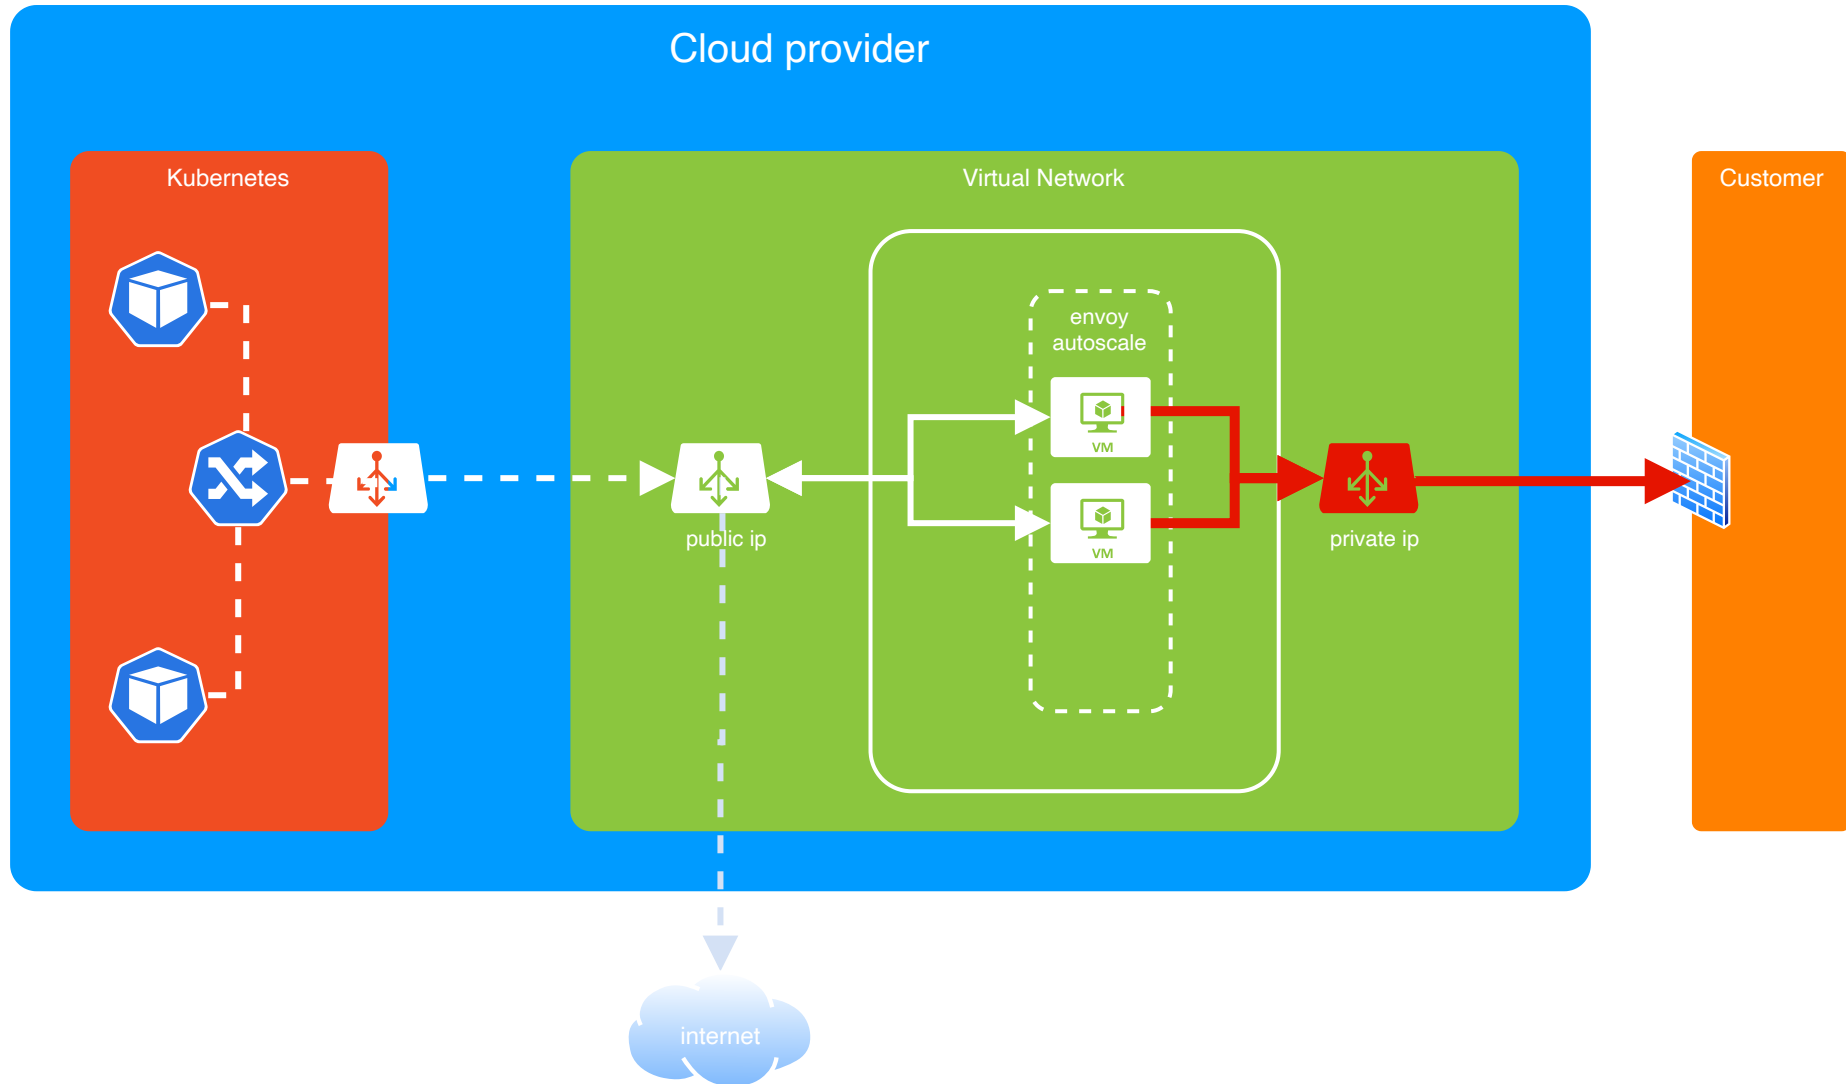




VNET can be privately connected to customer network

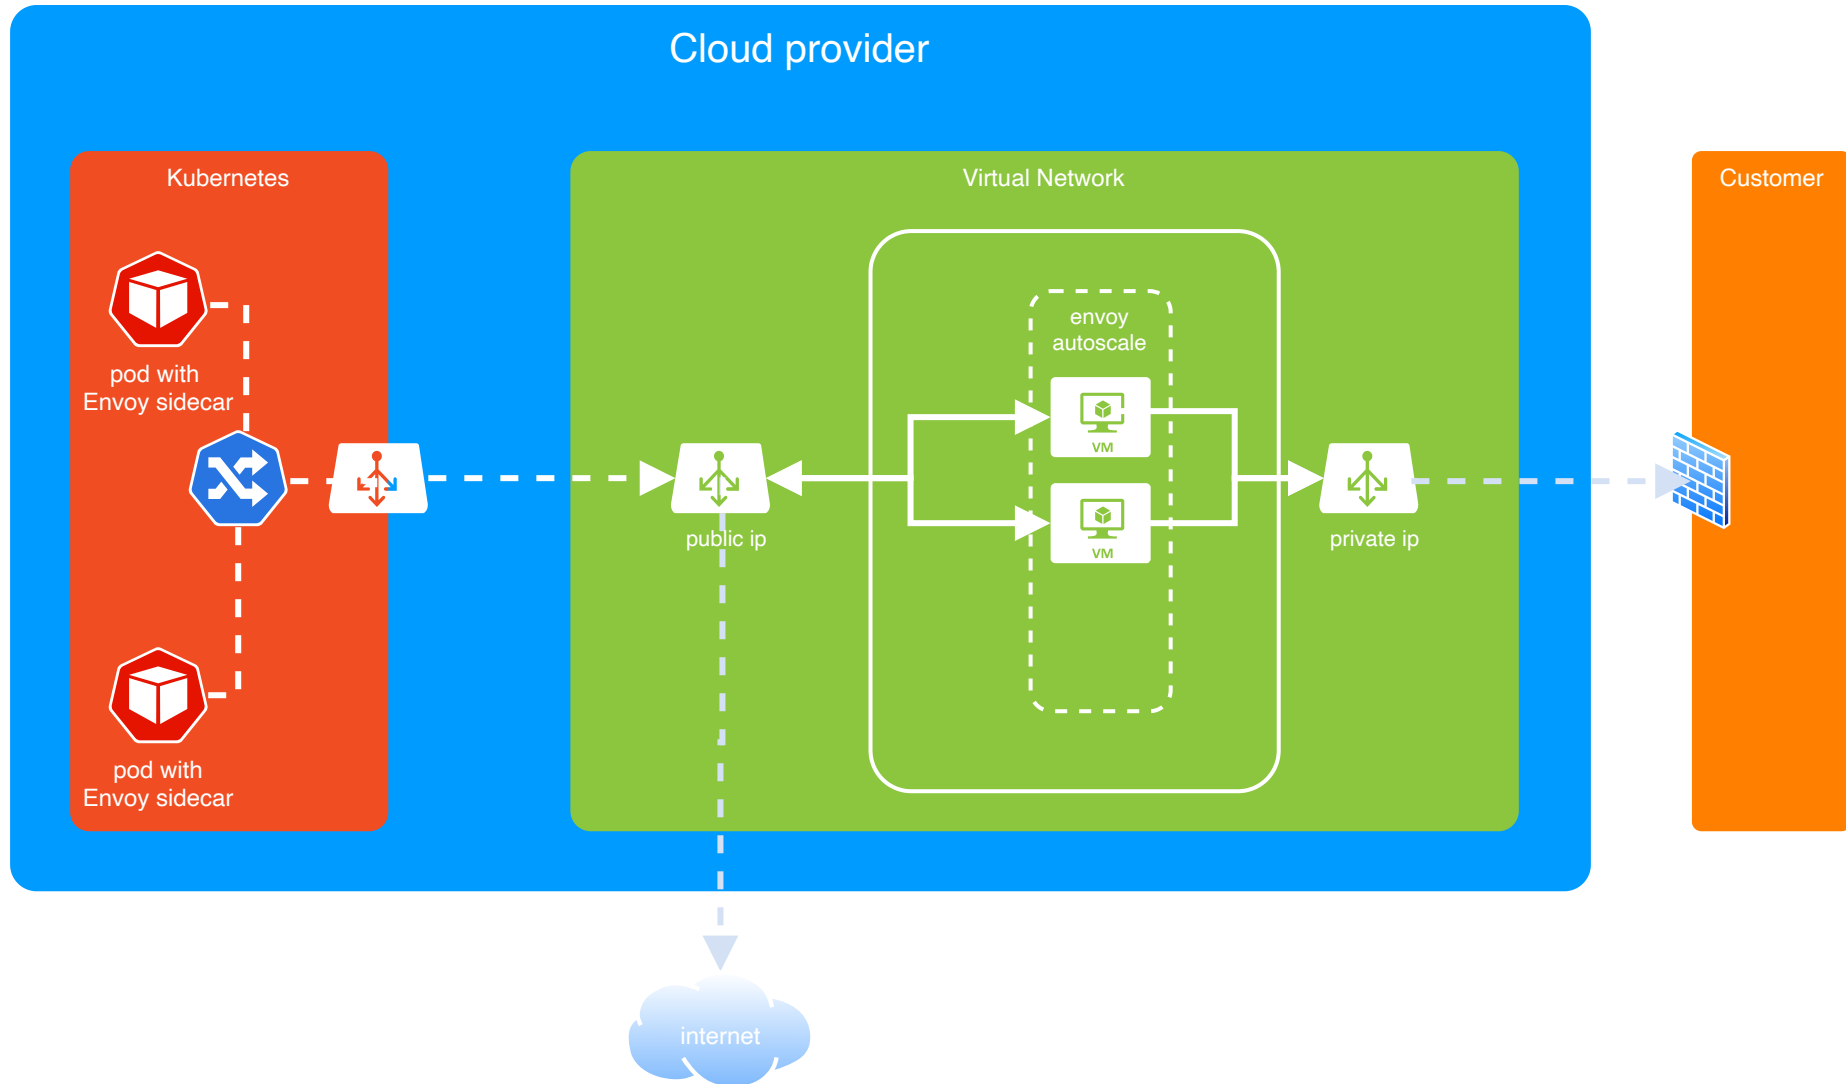


Load Balancer in front of VMs gives a dedicated public egress ip

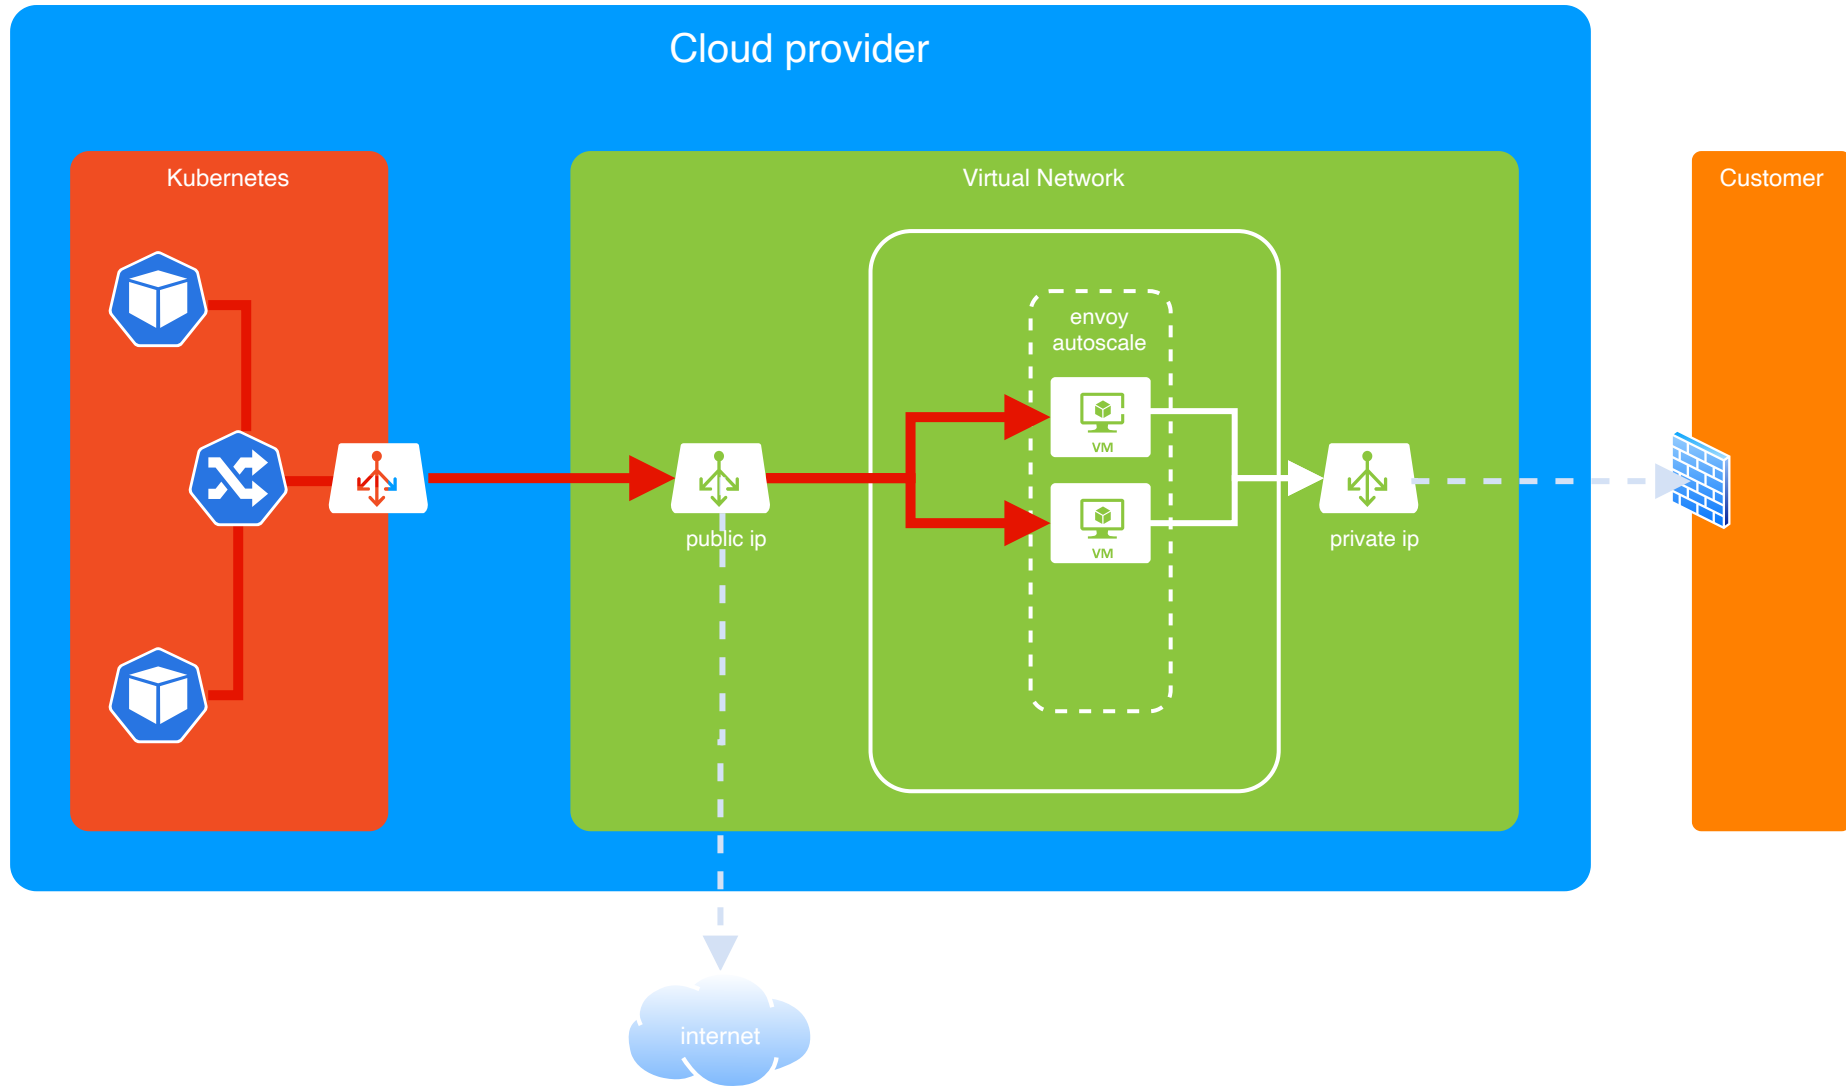


Private Load Balancer gives a dedicated private egress ip





JVM configured with Envoy sidecar as HTTP proxy



HTTP2 tunnel between sidecar Envoy and VM Envoy with mTLS

# ENVOY: PROS

Simple and transparent config in JVM using http proxy system properties

Any protocol supported using different listeners in sidecar

All traffic is encrypted

# ENVOY: PROS

VNET allows configuration of VPN, private connections at cloud level as a service

mTLS prevents unauthorized connections and one tenant to connect to another tenant Envoy

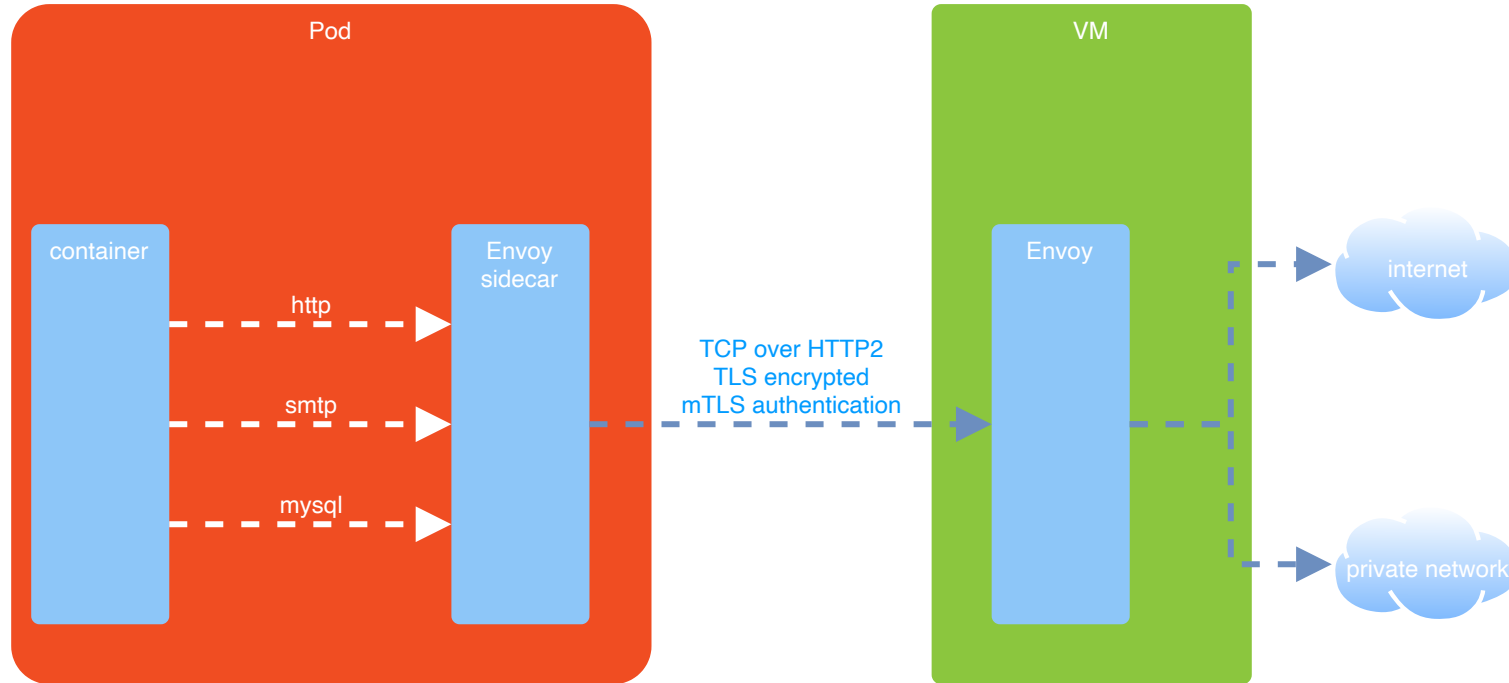
# ENVOY: CONS

Needs one set of certificates for each tenant for sidecars and  
VMs

rotation, expiration,...

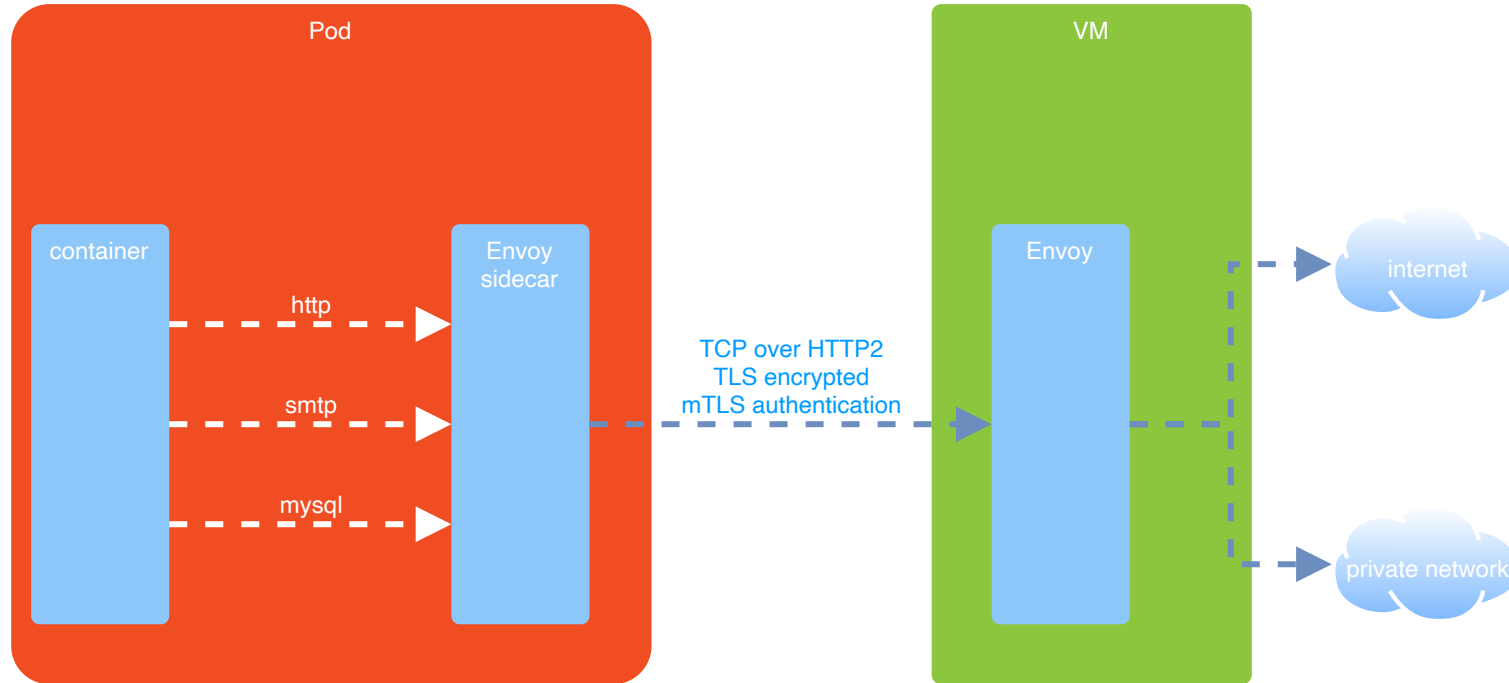
**ENVOY SIDECAR**

# ENVOY SIDECAR



One listener with TcpProxy filter for http/s. HTTP CONNECT gives Envoy the destination

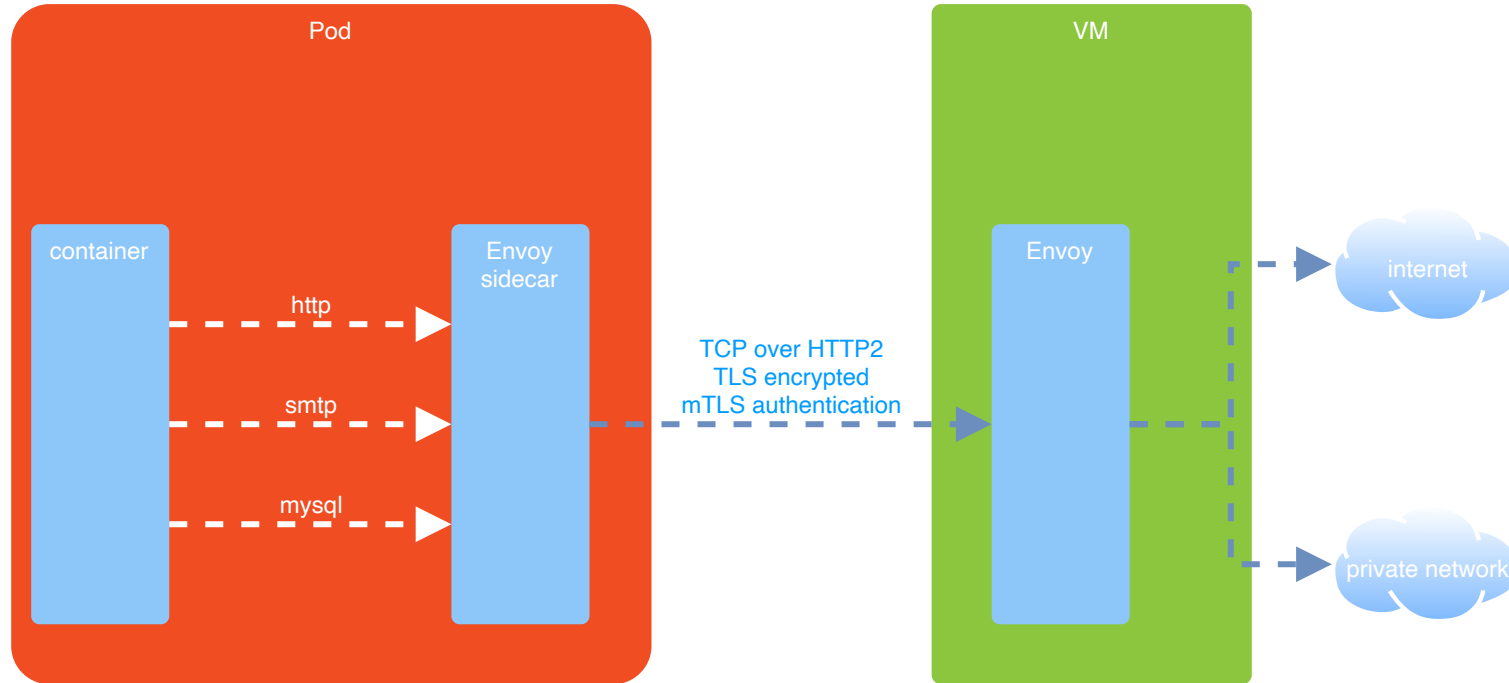
# ENVOY SIDECAR



One listener for each non http port. Destination hardcoded in  
`tunneling_config`



# ENVOY SIDECAR



One cluster with the VM Envoy LB as endpoint and TLS  
`transport_socket` config

# ENVOY SIDECAR: HTTP LISTENER

```
- name: listener_0
  "@type": type.googleapis.com/envoy.config.
listener.v3.Listener
  address:
    socket_address:
      protocol: TCP
      address: 0.0.0.0
      port_value: 3128
  filter_chains:
    - filters:
      - name: tcp
        typed_config:
          "@type": type.googleapis.com/
envoy.extensions
.filters.network.tcp_proxy.v3.TcpProxy
        stat_prefix: tcp_stats
        cluster: cluster_0
```

# ENVOY SIDECAR: NON HTTP LISTENER

```
- filters:  
  - name: tcp  
    typed_config:  
      "@type": type.googleapis.com/envoy.extensions.  
        filters.network.tcp_proxy.v3.TcpProxy  
      stat_prefix: tcp_stats  
      cluster: cluster_0  
      tunneling_config:  
        hostname: mysql:3306
```

# ENVOY SIDECAR: CLUSTER

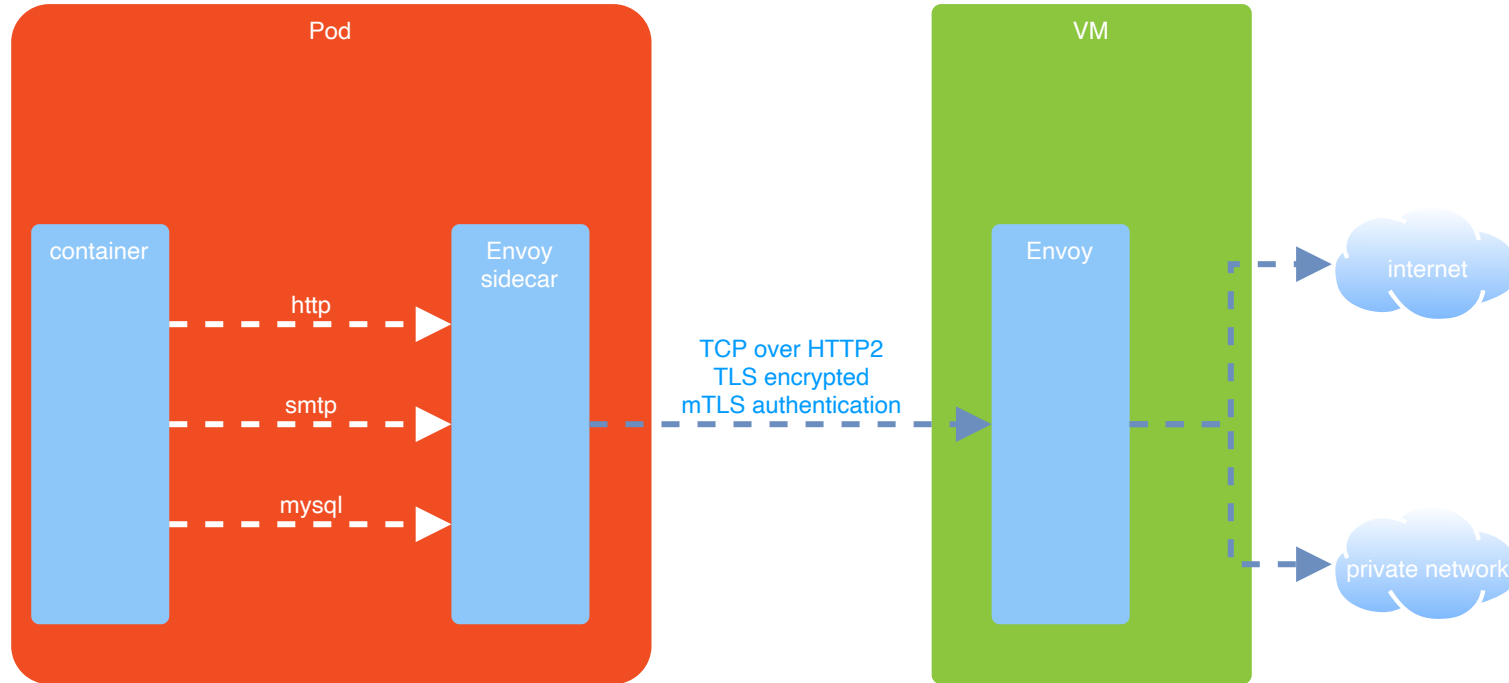
```
- name: cluster_0
  "@type": type.googleapis.com/envoy.
  config.cluster.v3.Cluster
  type: logical_dns
  typed_extension_protocol_options:
    envoy.extensions.upstreams.http.v3.
    HttpProtocolOptions:
      "@type": type.googleapis.com/envoy.extensions.
      upstreams.http.v3.HttpProtocolOptions
  respect_dns_ttl: true
  load_assignment:
    cluster_name: cluster_0
    endpoints:
      - lb_endpoints:
          - endpoint:
              address:
                socket_address:
                  address: envoy_vm
                  port_value: 443
```

# ENVOY SIDECAR: CLUSTER

```
transport_socket:
  name: envoy.transport_sockets.tls
  typed_config:
    "@type": type.googleapis.com/envoy.extensions.
    transport_sockets.tls.v3.UpstreamTlsContext
  common_tls_context:
    tls_certificate_sds_secret_configs:
      - name: upstream_cert
        sds_config:
          path: /var/lib/envoy/sds.yaml
    tls_params:
      tls_minimum_protocol_version: TLSv1_2
    validation_context_sds_secret_config:
      name: upstream_cert_ca
      sds_config:
        path: /var/lib/envoy/sds-ca.yaml
```

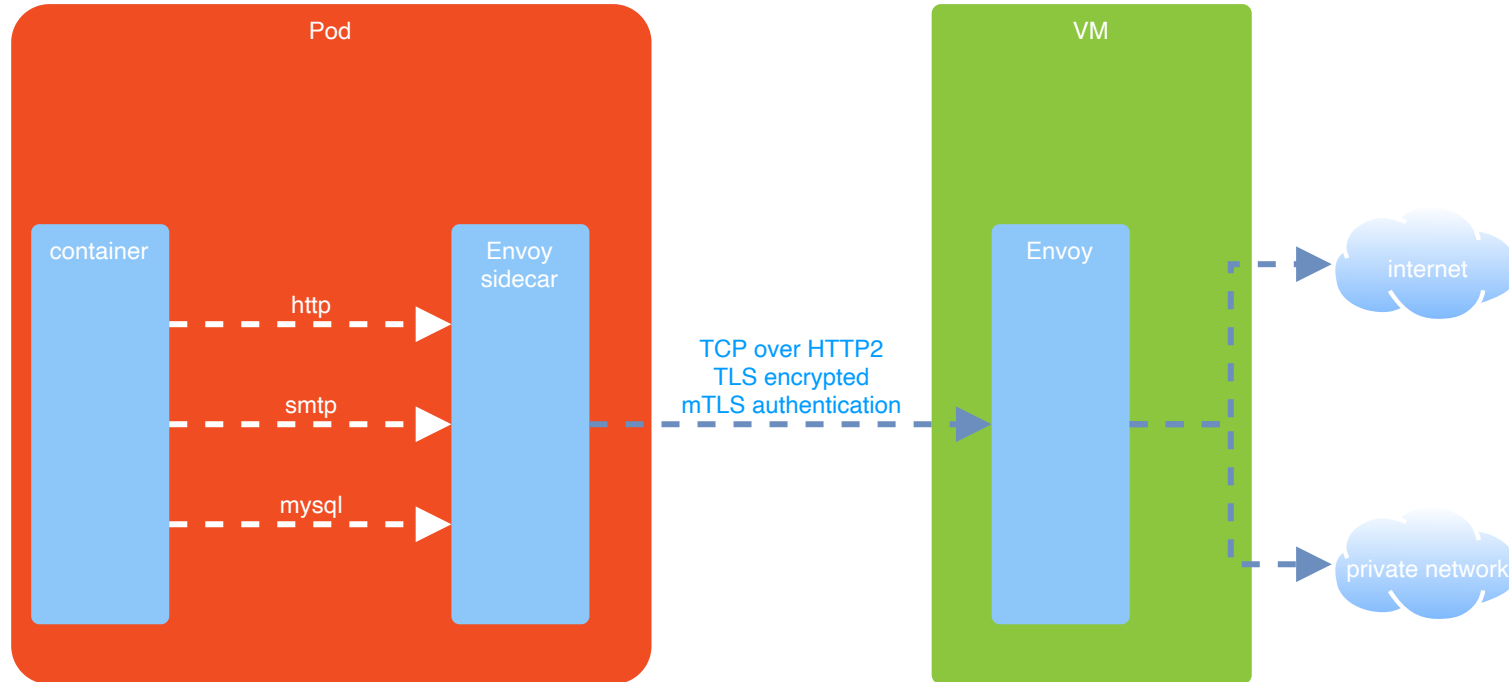
**ENVOY VM**

# ENVOY IN VM



One `HttpConnectionManager` listener with CONNECT upgrade

# ENVOY IN VM



One `dynamic_forward_proxy` cluster for all destinations



# ENVOY IN VM: LISTENER

```
- name: listener_0
  "@type": type.googleapis.com/envoy.config.listener.v3.Listener
  address:
    socket_address:
      protocol: TCP
      address: 0.0.0.0
      port_value: 443
  filter_chains:
    - filters:
        - name: envoy.filters.network.http_connection_manager
          typed_config:
            "@type": type.googleapis.com/envoy.extensions
              .filters.network.http_connection_manager.v3.
              HttpConnectionManager
```

# ENVOY IN VM: LISTENER

```
route_config:
  name: local_route
  virtual_hosts:
    - name: local_service
      domains: ["*"]
      routes:
        - match:
            connect_matcher: {}
          route:
            cluster: dynamic_forward_proxy_cluster
            upgrade_configs:
              - upgrade_type: CONNECT
                connect_config: {}
        # needed to be used as a proxy with http (not s)
        - match:
            prefix: "/"
          route:
            cluster: dynamic_forward_proxy_cluster
```

# ENVOY IN VM: LISTENER

```
http_filters:
- name: envoy.filters.http.dynamic_forward_proxy
  typed_config:
    "@type": type.googleapis.com/envoy.extensions.
      filters.http.dynamic_forward_proxy.v3.FilterConfig
    dns_cache_config:
      name: dynamic_forward_proxy_cache_config
      dns_lookup_family: V4_ONLY
- name: envoy.filters.http.router
  typed_config:
    "@type": type.googleapis.com/envoy.extensions.
      filters.http.router.v3.Router
upgrade_configs:
- upgrade_type: CONNECT
```

# ENVOY IN VM: LISTENER

```
transport_socket:
  name: envoy.transport_sockets.tls
  typed_config:
    "@type": type.googleapis.com/envoy.extensions.
    transport_sockets.tls.v3.DownstreamTlsContext
  common_tls_context:
    tls_certificate_sds_secret_configs:
      - name: upstream_cert
        sds_config:
          path: /var/lib/envoy/sds.yaml
    tls_params:
      tls_minimum_protocol_version: TLSv1_2
    validation_context_sds_secret_config:
      name: upstream_cert_ca
      sds_config:
        path: /var/lib/envoy/sds-ca.yaml
  require_client_certificate: true
```

# ENVOY IN VM: CLUSTER

```
- name: dynamic_forward_proxy_cluster
  "@type": type.googleapis.com/envoy.config.cluster.v3.Cluster
  connect_timeout: 30s
  lb_policy: CLUSTER_PROVIDED
  cluster_type:
    name: envoy.clusters.dynamic_forward_proxy
    typed_config:
      "@type": type.googleapis.com/envoy.extensions.clusters.dynamic_forward_proxy.v3.ClusterConfig
      dns_cache_config:
        name: dynamic_forward_proxy_cache_config
        dns_lookup_family: V4_ONLY
```

# ENVOY IN VM: SDS

```
resources:  
  - "@type": type.googleapis.com/envoy.extensions.  
    transport_sockets.tls.v3.Secret  
    name: upstream_cert_ca  
    validation_context:  
      trusted_ca:  
        filename: /etc/envoy/certs/cacert.pem  
        # only allow connections with certificates  
        # that have this SAN  
      match_subject_alt_names:  
        - exact: "envoy_sidecar"
```

**CERTIFICATE ROTATION**

# CERTIFICATE ROTATION

Using Vault Agent to generate short lived certificates (sidecar and VM)

CA stored in Vault

Kubernetes and Azure passwordless authentication to Vault



# ENVOY SIDECAR: SDS

Configuring certificates as SDS in separate files to ensure certificates are automatically reloaded on change.

```
resources:  
  - "@type": type.googleapis.com/envoy.extensions.  
    transport_sockets.tls.v3.Secret  
    name: upstream_cert  
    tls_certificate:  
      certificate_chain:  
        filename: "/etc/envoy/certs/tls.crt"  
      private_key:  
        filename: "/etc/envoy/certs/tls.key"
```

# ENVOY SIDECAR: SDS-CA

```
resources:  
  - "@type": type.googleapis.com/envoy.extensions.  
    transport_sockets.tls.v3.Secret  
    name: upstream_cert_ca  
    validation_context:  
      trusted_ca:  
        filename: "/etc/envoy/certs/cacert.pem"
```

# **CERTIFICATE ROTATION: ALTERNATIVES**

Can use certmanager to do the rotation

only in K8S

Spiffe is another option

# ENVOY DEBUGGING

# ENVOY DEBUGGING

TLS connection errors only show up in `connection`  
component debug logs

Client only sees socket closing messages

# OPENSSL\_internal ERRORS

KEY\_VALUES\_MISMATCH [warning/config]

key does not match cert

SSLV3\_ALERT\_CERTIFICATE\_EXPIRED [debug/connection]

CERTIFICATE\_VERIFY\_FAILED

SSLV3\_ALERT\_CERTIFICATE\_UNKNOWN [debug][connection]

cert SAN does not match expected SAN

# Example: certificate SAN does not match match\_subject\_alt\_names

## VM side

```
envoy_vm_1      [debug][connection]  
[source/extensions/transport_sockets/tls/ssl_socket.cc:224]  
[C0] TLS error: 268435581:SSL routines:  
OPENSSL_internal:CERTIFICATE_VERIFY_FAILED
```

## Sidecar side

```
envoy_sidecar_1 [debug][connection]  
[source/extensions/transport_sockets/tls/ssl_socket.cc:224]  
[C1] TLS error: 268436502:SSL routines:  
OPENSSL_internal:SSLV3_ALERT_CERTIFICATE_UNKNOWN  
envoy_sidecar_1 [debug][connection]  
[source/common/network/connection_impl.cc:241]  
[C1] closing socket: 0
```

# HOW IS IT USED

- http/s: as a http proxy
- other ports: open port in localhost



# HTTP PROXY: JAVA

## System properties

- `http.proxyHost` & `http.proxyPort`
- `https.proxyHost` & `https.proxyPort`
- `http.nonProxyHosts`

# HTTP PROXY: JAVA

Apache HTTP client ignores proxy by default

```
HttpClient httpClient = HttpClientBuilder.create().  
    useSystemProperties().build();
```

# LONG RUNNING CONNECTIONS

Connection pools or long running connections are typical (ie. db)

Configure `max_stream_duration` and  
`stream_idle_timeout`

Increase load balancer timeouts

# CONNECTION POOLS

Must configure connection pools to validate connections before  
use

In JDBC drivers it is a configuration option

# ENVOY TWEAKS

Some defaults can to be tweaked on the VM side listener

```
typed_config:
  "@type": type.googleapis.com/envoy.extensions
    .filters.network.http_connection_manager.v3.
    HttpConnectionManager
  common_http_protocol_options:
    # do not reset connections if they are not idle, set to 24h
    max_stream_duration: 86400s
  # increase the stream idle timeout, it ignores tcp keepalives
  stream_idle_timeout: 7200s
```

# HTTP PROXY: HTTPD

```
SSLProxyEngine on

ProxyRemote "https://example.com"
            "http://${HTTP_PROXY_HOST}:${HTTP_PROXY_PORT}"
ProxyPass "/somepath" "https://example.com"
ProxyPassReverse "/somepath" "https://example.com"
```

httpd sends proxy requests with CONNECT HTTP/1.0

we need to set http\_protocol\_options.

accept\_http\_10

# ENVOY CONFIG: RESOURCES

envoyproxy.io

[arch\\_overview/http/upgrades](#)

[sandboxes/tls](#)

[sandboxes/double-proxy](#)



[csanchez.org](http://csanchez.org)

